

CS 475 Machine Learning: Homework 4 Analytical (70 points)

Assigned: Monday, Nov. 1st, 2021

Due: Monday, Nov. 15th, 2021, 11:59 pm US/Eastern

Partner 1: NAME (JHED), Partner 2: NAME (JHED)

Instructions

We have provided this L^AT_EX document for turning in this homework. We give you one or more boxes to answer each question. The question to answer for each box will be noted in the title of the box. You can change the size of the box if you need more space.

Other than your name, do not type anything outside the boxes. Leave the rest of the document unchanged.

Do not add text outside of the answer boxes. You are allowed to make boxes larger if needed.

We strongly recommend you review your answers in the generated PDF to ensure they appear correct. We will grade what appears in the answer boxes in the submitted PDF, NOT the original latex file.

MRFs

Question 1.

Consider the graphical model shown in Figure 1. In this model, \mathbf{x} is a sequence of observations for which we want to output a prediction \mathbf{y} , which itself is a sequence, where the size of \mathbf{y} is the same as \mathbf{x} . Assume that the potential functions have a log-linear form: $\psi(Z) = \exp\{\sum_i \theta_i f_i(Z)\}$, where Z is the set of nodes that are arguments to the potential function (i.e. some combination of nodes in \mathbf{x} and \mathbf{y}), θ are the parameters of the potential functions and f_i is a feature function.




Figure 1: Tree structure model

- Write the log likelihood for this model of a single instance \mathbf{x} : $\log p(\mathbf{y}, \mathbf{x})$.
- Write the conditional log likelihood for this model of a single instance \mathbf{x} : $\log p(\mathbf{y}|\mathbf{x})$.
- Assume that each variable y_i can take one of k possible states, and variable x_i can take one of k' possible states, where k' is very large. Describe the computational challenges of modeling $\log p(\mathbf{y}, \mathbf{x})$ vs $\log p(\mathbf{y}|\mathbf{x})$.

Question 2.


- (a) Suppose you wanted to compute $S = \sum_{x_1=1}^{100} \cdots \sum_{x_8=1}^{100} h(x)$ where

$$h(x) = \exp(x_1x_2 + x_4x_5 + x_7x_8) \prod_{i=2,5,7} (x_i + x_3 + x_6)^i.$$

It looks like the sum has $100^8 = 10^{16}$ terms, so it seems we must evaluate h 10^{16} times. Explain (precisely) how you can compute S with at most 10^7 evaluations of h or something simpler than h .

- (b) Draw the MRF associated with this distribution.

DAGs, Clique Trees and Message Passing.



In a statistical DAG model for the graph shown, let $\mathbf{V} = \{A, B, C, D, E, F, G, H\}$.

- (a) Answer (and explain your answer) the following d-separation queries:

$$A \perp\!\!\!\perp F \mid D$$

$$A \perp\!\!\!\perp G \mid B, C$$

$$G \perp\!\!\!\perp A \mid B, H, D, E, F$$

$$F \perp\!\!\!\perp D \mid A, B$$

$$C \perp\!\!\!\perp H \mid B$$

- (b) Write down the local Markov property of this model.
- (c) Consider a new graph where we reverse the direction of the edge $B \rightarrow G$ to point the other way: $B \leftarrow G$ (and leave the other edges the same). Does the new graph represent the same model as the old?

Hint: write down the local Markov property for the new graph, and see if all statements in it are implied by d-separation in the original graph. In general, if local Markov of \mathcal{G}_1 is implied by global Markov of \mathcal{G}_2 , and local Markov of \mathcal{G}_2 is implied by global Markov of \mathcal{G}_1 , then \mathcal{G}_1 and \mathcal{G}_2 represent the same model. Otherwise they do not.

- (d) A moralized graph \mathcal{G}^a is obtained from a DAG \mathcal{G} by connecting all non-adjacent variables V_i and V_j such that $V_i \rightarrow V_k \leftarrow V_j$ is in the graph (for some V_k), and replacing all directed edges by undirected edges. What is the moralized graph for the DAG in this problem?
- (e) Write down the MRF factorization of the moralized graph \mathcal{G}^a .
- (f) Is this graph chordal? If not, add a set of edges to make it chordal. If you added edges, write the factorization of the new graph.
- (g) Create a clique tree from the triangulated graph (either \mathcal{G}^a or the graph obtained from \mathcal{G}^a by adding new edge(s)).
- (h) Pick a root \mathbf{R} of the clique tree, and calculate both incoming messages $\phi^{\mathbf{S}_i \rightarrow \mathbf{S}_j}$ from each \mathbf{S}_i towards its neighbor \mathbf{S}_j closer to the root, and outgoing messages $\phi^{\mathbf{S}_k \leftarrow \mathbf{S}_i}$ from \mathbf{S}_i to each neighbor \mathbf{S}_k further than \mathbf{S}_i from the root, in terms of clique potentials and other messages.
- (i) By substituting in the clique factors in each message, show that in this example, for each leaf node \mathbf{S}_i with a neighbor node \mathbf{S}_j ,

$$p(\mathbf{S}_i) = \frac{\phi^{\mathbf{S}_i \leftarrow \mathbf{S}_j} \phi_{\mathbf{S}_i}}{\sum_{\mathbf{S}_i} \phi^{\mathbf{S}_i \leftarrow \mathbf{S}_j} \phi_{\mathbf{S}_i}} = \frac{\sum_{\mathbf{V} \setminus \mathbf{S}_i} \prod_{C \in \mathcal{C}(\mathcal{G})} \phi_C}{\sum_{\mathbf{V}} \prod_{C \in \mathcal{C}(\mathcal{G})} \phi_C}$$

for each non-leaf node \mathbf{S}_i with a neighbor \mathbf{S}_j closer to the root, and neighbors $\mathbf{S}_1, \dots, \mathbf{S}_m$ further from the root that

$$p(\mathbf{S}_i) = \frac{\phi_{\mathbf{S}_j \setminus \mathbf{S}_i}^{\mathbf{S}_i \leftarrow \mathbf{S}_j} \left(\prod_{k=1}^m \phi_{\mathbf{S}_k \setminus \mathbf{S}_i}^{\mathbf{S}_k \rightarrow \mathbf{S}_i} \right) \phi_{\mathbf{S}_i}}{\sum_{\mathbf{S}_i} \phi_{\mathbf{S}_j \setminus \mathbf{S}_i}^{\mathbf{S}_i \leftarrow \mathbf{S}_j} \left(\prod_{k=1}^m \phi_{\mathbf{S}_k \setminus \mathbf{S}_i}^{\mathbf{S}_k \rightarrow \mathbf{S}_i} \right) \phi_{\mathbf{S}_i}} = \frac{\sum_{\mathbf{V} \setminus \mathbf{S}_i} \prod_{C \in \mathcal{C}(\mathcal{G})} \phi_C}{\sum_{\mathbf{V}} \prod_{C \in \mathcal{C}(\mathcal{G})} \phi_C}$$

and finally for the root node \mathbf{S}_i with neighbors $\mathbf{S}_1, \dots, \mathbf{S}_m$ that

$$p(\mathbf{S}_i) = \frac{\left(\prod_{k=1}^m \phi_{\mathbf{S}_k \setminus \mathbf{S}_i}^{\mathbf{S}_k \rightarrow \mathbf{S}_i} \right) \phi_{\mathbf{S}_i}}{\sum_{\mathbf{S}_i} \left(\prod_{k=1}^m \phi_{\mathbf{S}_k \setminus \mathbf{S}_i}^{\mathbf{S}_k \rightarrow \mathbf{S}_i} \right) \phi_{\mathbf{S}_i}} = \frac{\sum_{\mathbf{V} \setminus \mathbf{S}_i} \prod_{C \in \mathcal{C}(\mathcal{G})} \phi_C}{\sum_{\mathbf{V}} \prod_{C \in \mathcal{C}(\mathcal{G})} \phi_C}$$

Here \mathbf{V} is all variables in the graph, and $\mathcal{C}(\mathcal{G})$ is the set of maximal cliques in the graph.

K-Means

- (a) Is it possible to initialize the k-means algorithm in such a way that it fails to terminate successfully?
- (b) Say our input to k-means is a set of $2k$ points with 2-coordinates arranged in line, e.g. with coordinates:

$$(0, 0), (0, 1), (0, 2), \dots, (0, k), (0, k+1), \dots, (0, 2k).$$

Say we initialize k -means with 2 clusters, with initial centroids given by $(0, k)$ and $(0, k+1)$. In many iterations will k -means terminate? What will be the final cluster assignments and centroids?